

# Apache Spark and Apache Flink Comparison

---

**Content- Mangement und Suchtechnologien**

Florian Thom s0558101

Daniel Nagel s0559090

# Apache Spark: Streaming

---

»Spark Streaming is an extension of the core Spark API that enables scalable, high-throughput, fault-tolerant stream processing of live data streams.«

- Spark Official Documentation

- **descretized stream (DStream) represents a continuous stream of data**
  - internally represented by a series of RDDs
- **DStream provides higher-order functions**
  - all operations are working on the underlying RDDs
- **two built-in streaming sources: basic and advanced**
  - basic: file systems and socket connections
  - advanced: Apache Kafka, Apache Flume, Amazon Kinesis

# Apache Spark: DStream



- **DStream is split into batches with RDDs as data representation**
  - multiple batches can be stored in a single window
- **StreamingContext contains the execution threads and batch interval**
  - the batch interval defines when current batch ends and a new batch starts
  - batches can be grouped into windows
  - window interval needs to be a multiple of the batch interval
- **use foreachRDD to access the underlying RDD representation**

# Apache Spark: Streaming Example

---



```
val sc = spark.sparkContext
val ssc = new StreamingContext(sc, Seconds(2))

var predicted: RDD[(Vector, Int)] = sc.emptyRDD
var kmodel = new KMeansModel(Array.fill(numClusters)(Vectors.zeros(numDim)))
val data = ssc.textFileStream("file://" + relPath.getCanonicalPath)
    .map(parseCSV(_))
    .filter(!containsEmpty(_))
    .map(transformToTuple(_))
val vectors = data.map(f => Vectors.dense(f._4, f._5, f._7))

vectors.foreachRDD(rdd => {
    kmodel = new KMeans().setK(numClusters).setInitialModel(kmodel).setMaxIterations(numIter).run(rdd)
    val current = rdd.map(f => (f, kmodel.predict(f)))
    predicted = predicted.union(current)
})

ssc.start()
ssc.awaitTerminationOrTimeout(20000)
ssc.stop(false)
```

# Apache Spark: Streaming Example




- exoplanet archive dataset from caltech and nasa
- clustering results can be exported at runtime

# Apache Flink: Streaming I

»Apache Flink is a framework and distributed processing engine for stateful computations over unbounded and bounded data streams.«

- Flink Official



- **bounded datastreams (near real-time, similar to Apache Spark)**
  - have a defined start and end
- **unbounded stream (real-time)**
  - have a start but no defined end
- **stateful functions and operators store data across the process events**

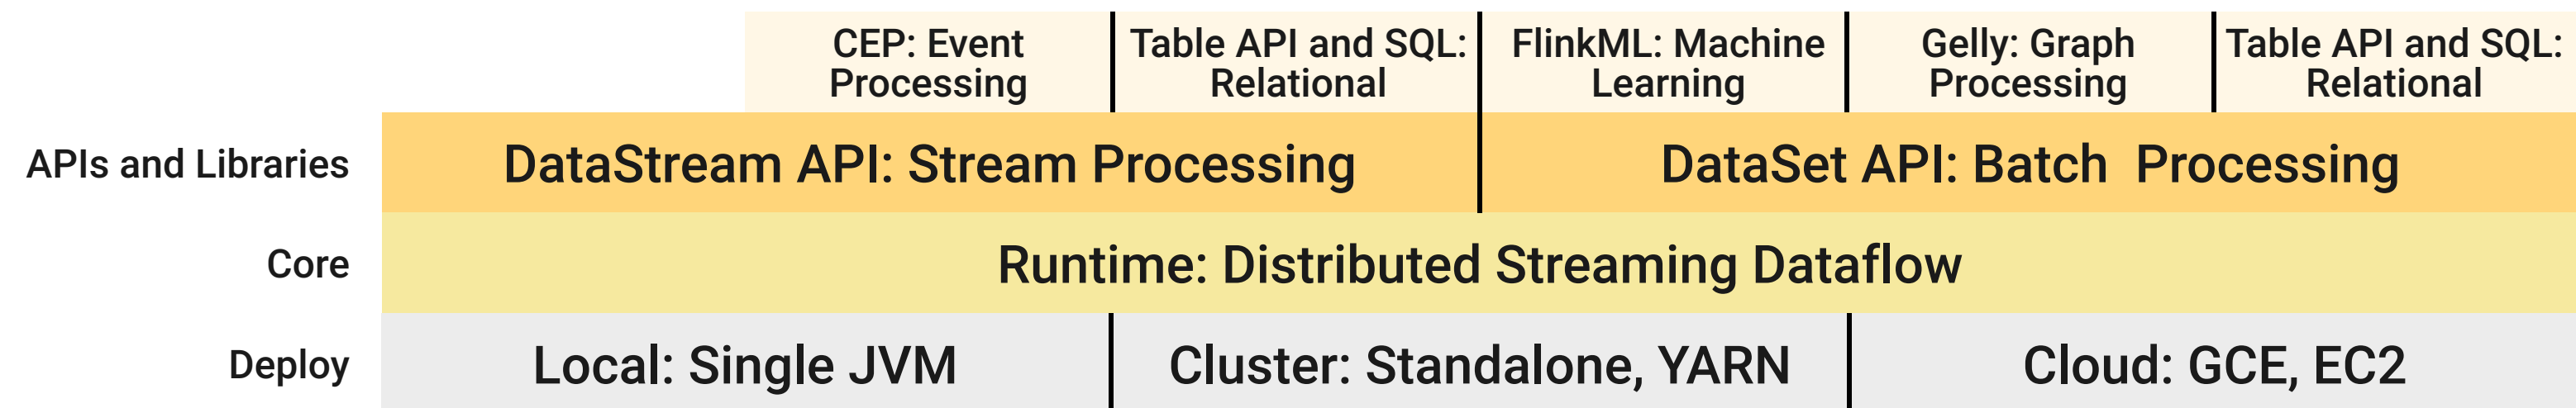
# Apache Flink: Streaming II

---

- **high throughput, scalable**
- **low latency due real time computation**
  - no waiting required (window / batch)
- **exactly-once semantics for stateful computations**
  - Flink: Standard
  - Spark: only with much efforts, not applicable to windowed operations
- **base implementation in Java**



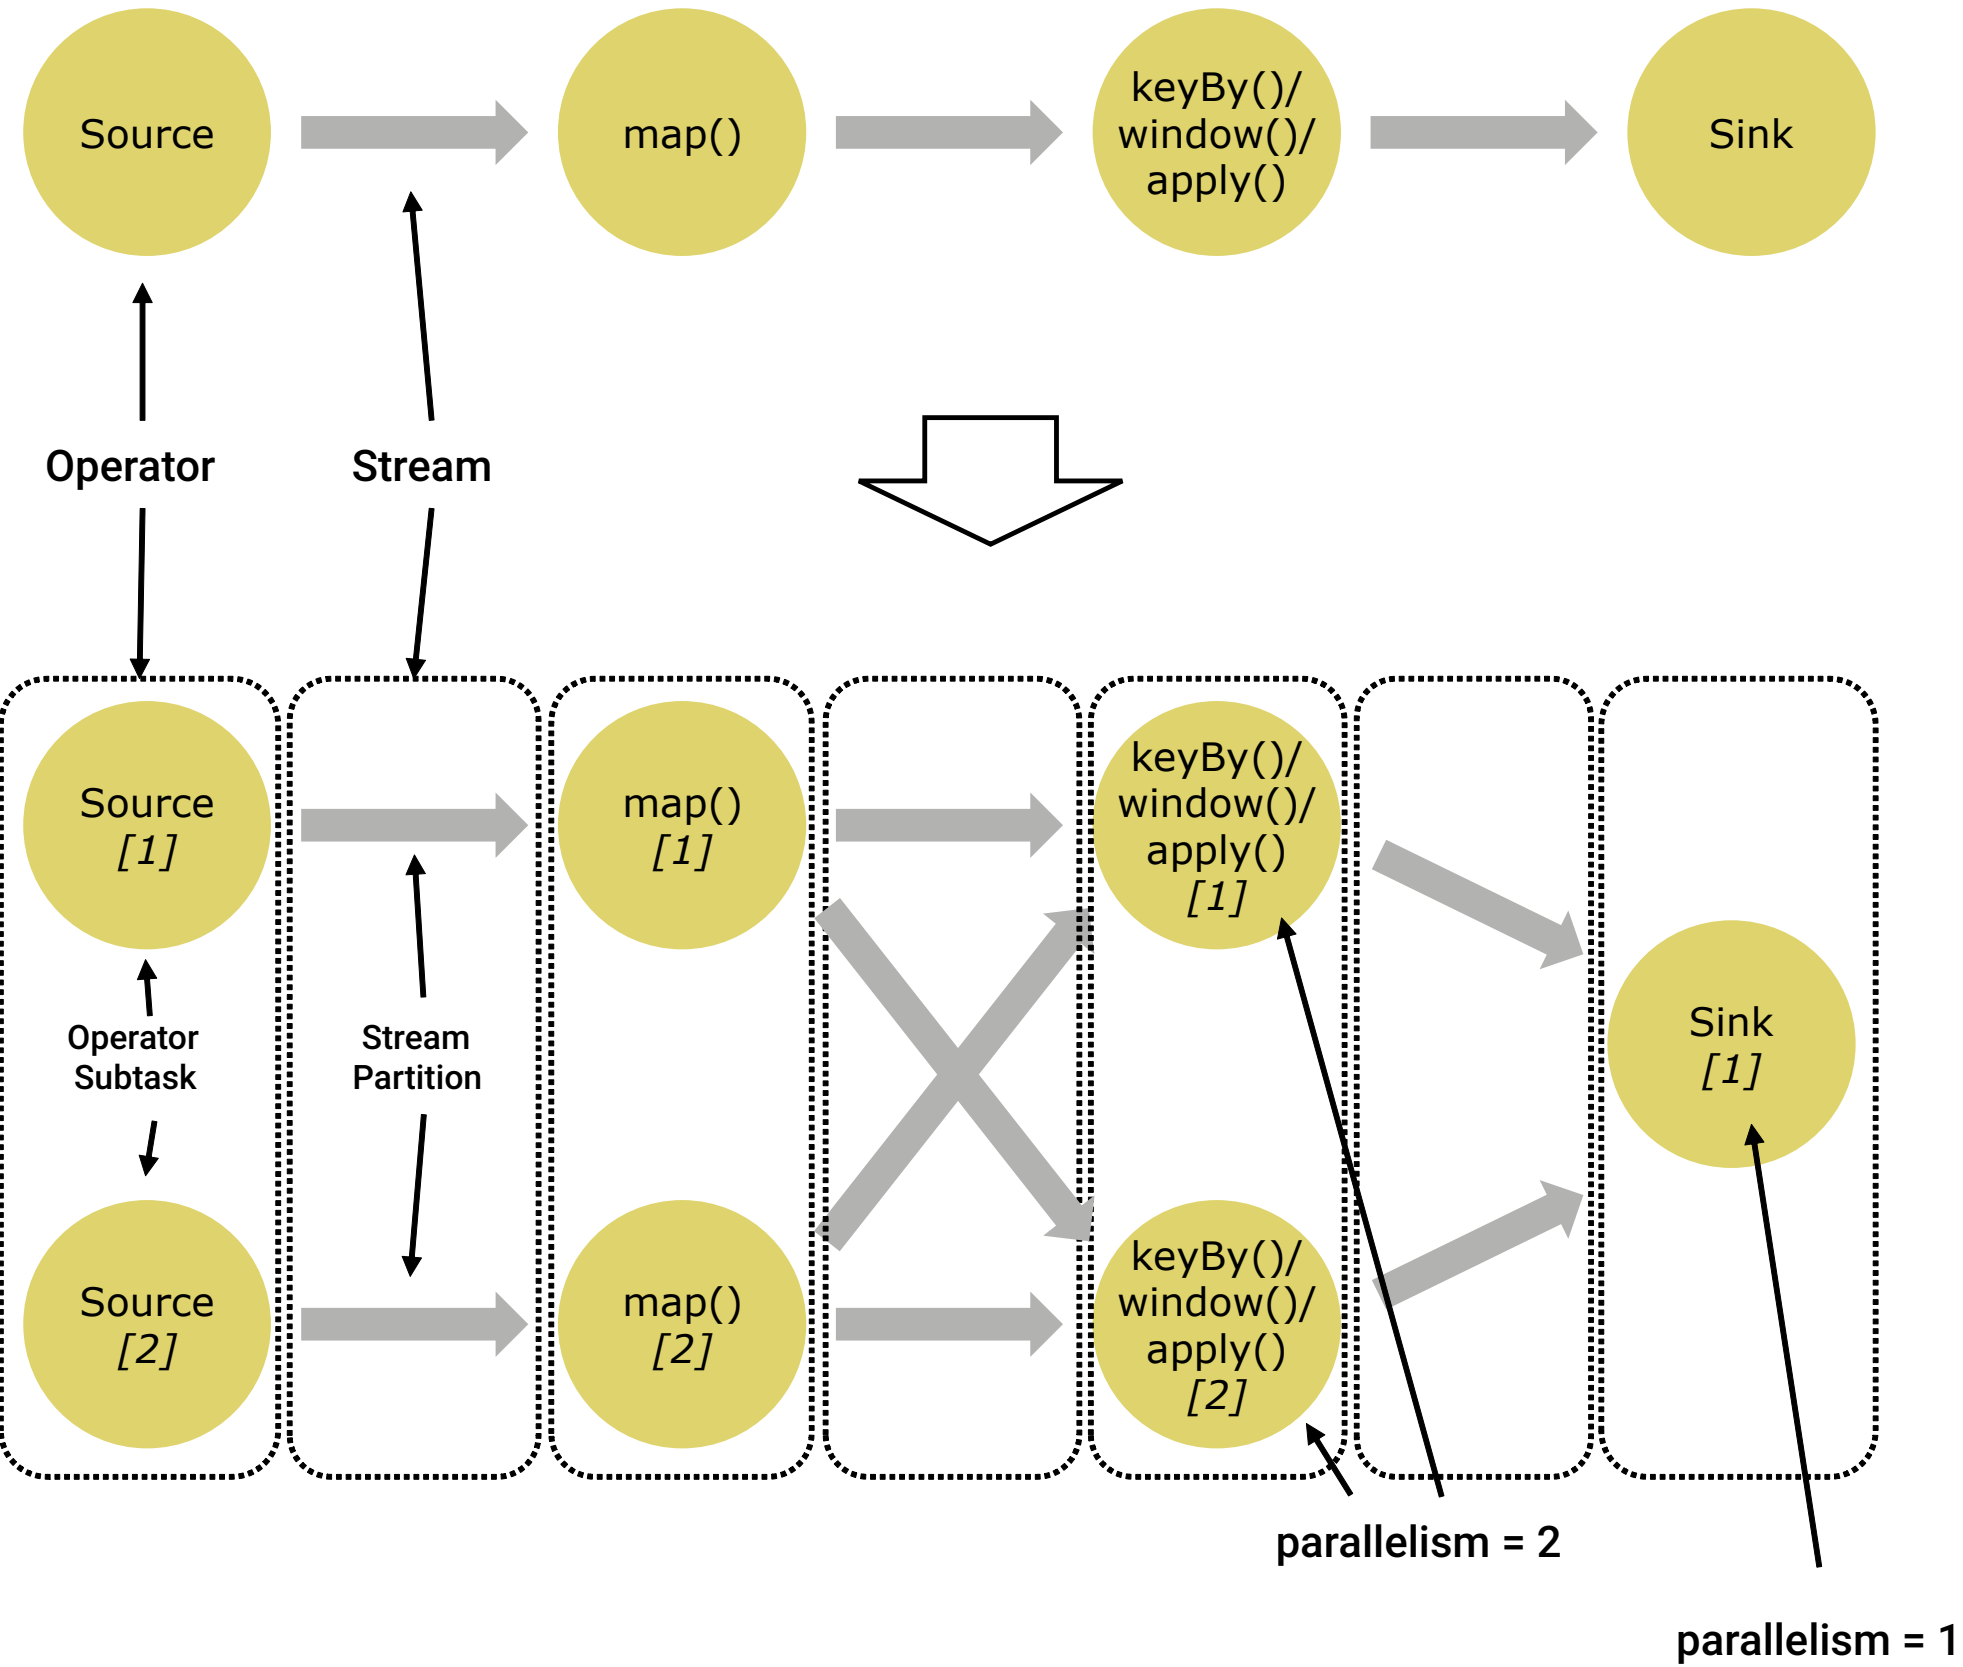
# Apache Flink: API I



- **DataSet: base data structure for batch processing**
  - internally represented as bounded stream
- **DataStream: base data structure for stream processing**
  - internal representation as directed acyclic graph (job graph)
  - start: data source / connectors (build-in: Apache Kafka, RabbitMQ, ...)
  - end: data sink / connectors (build-in: Apache Kafka, Cassandra, Redis, ...)
- **DataSet and DataStream can not be combined**

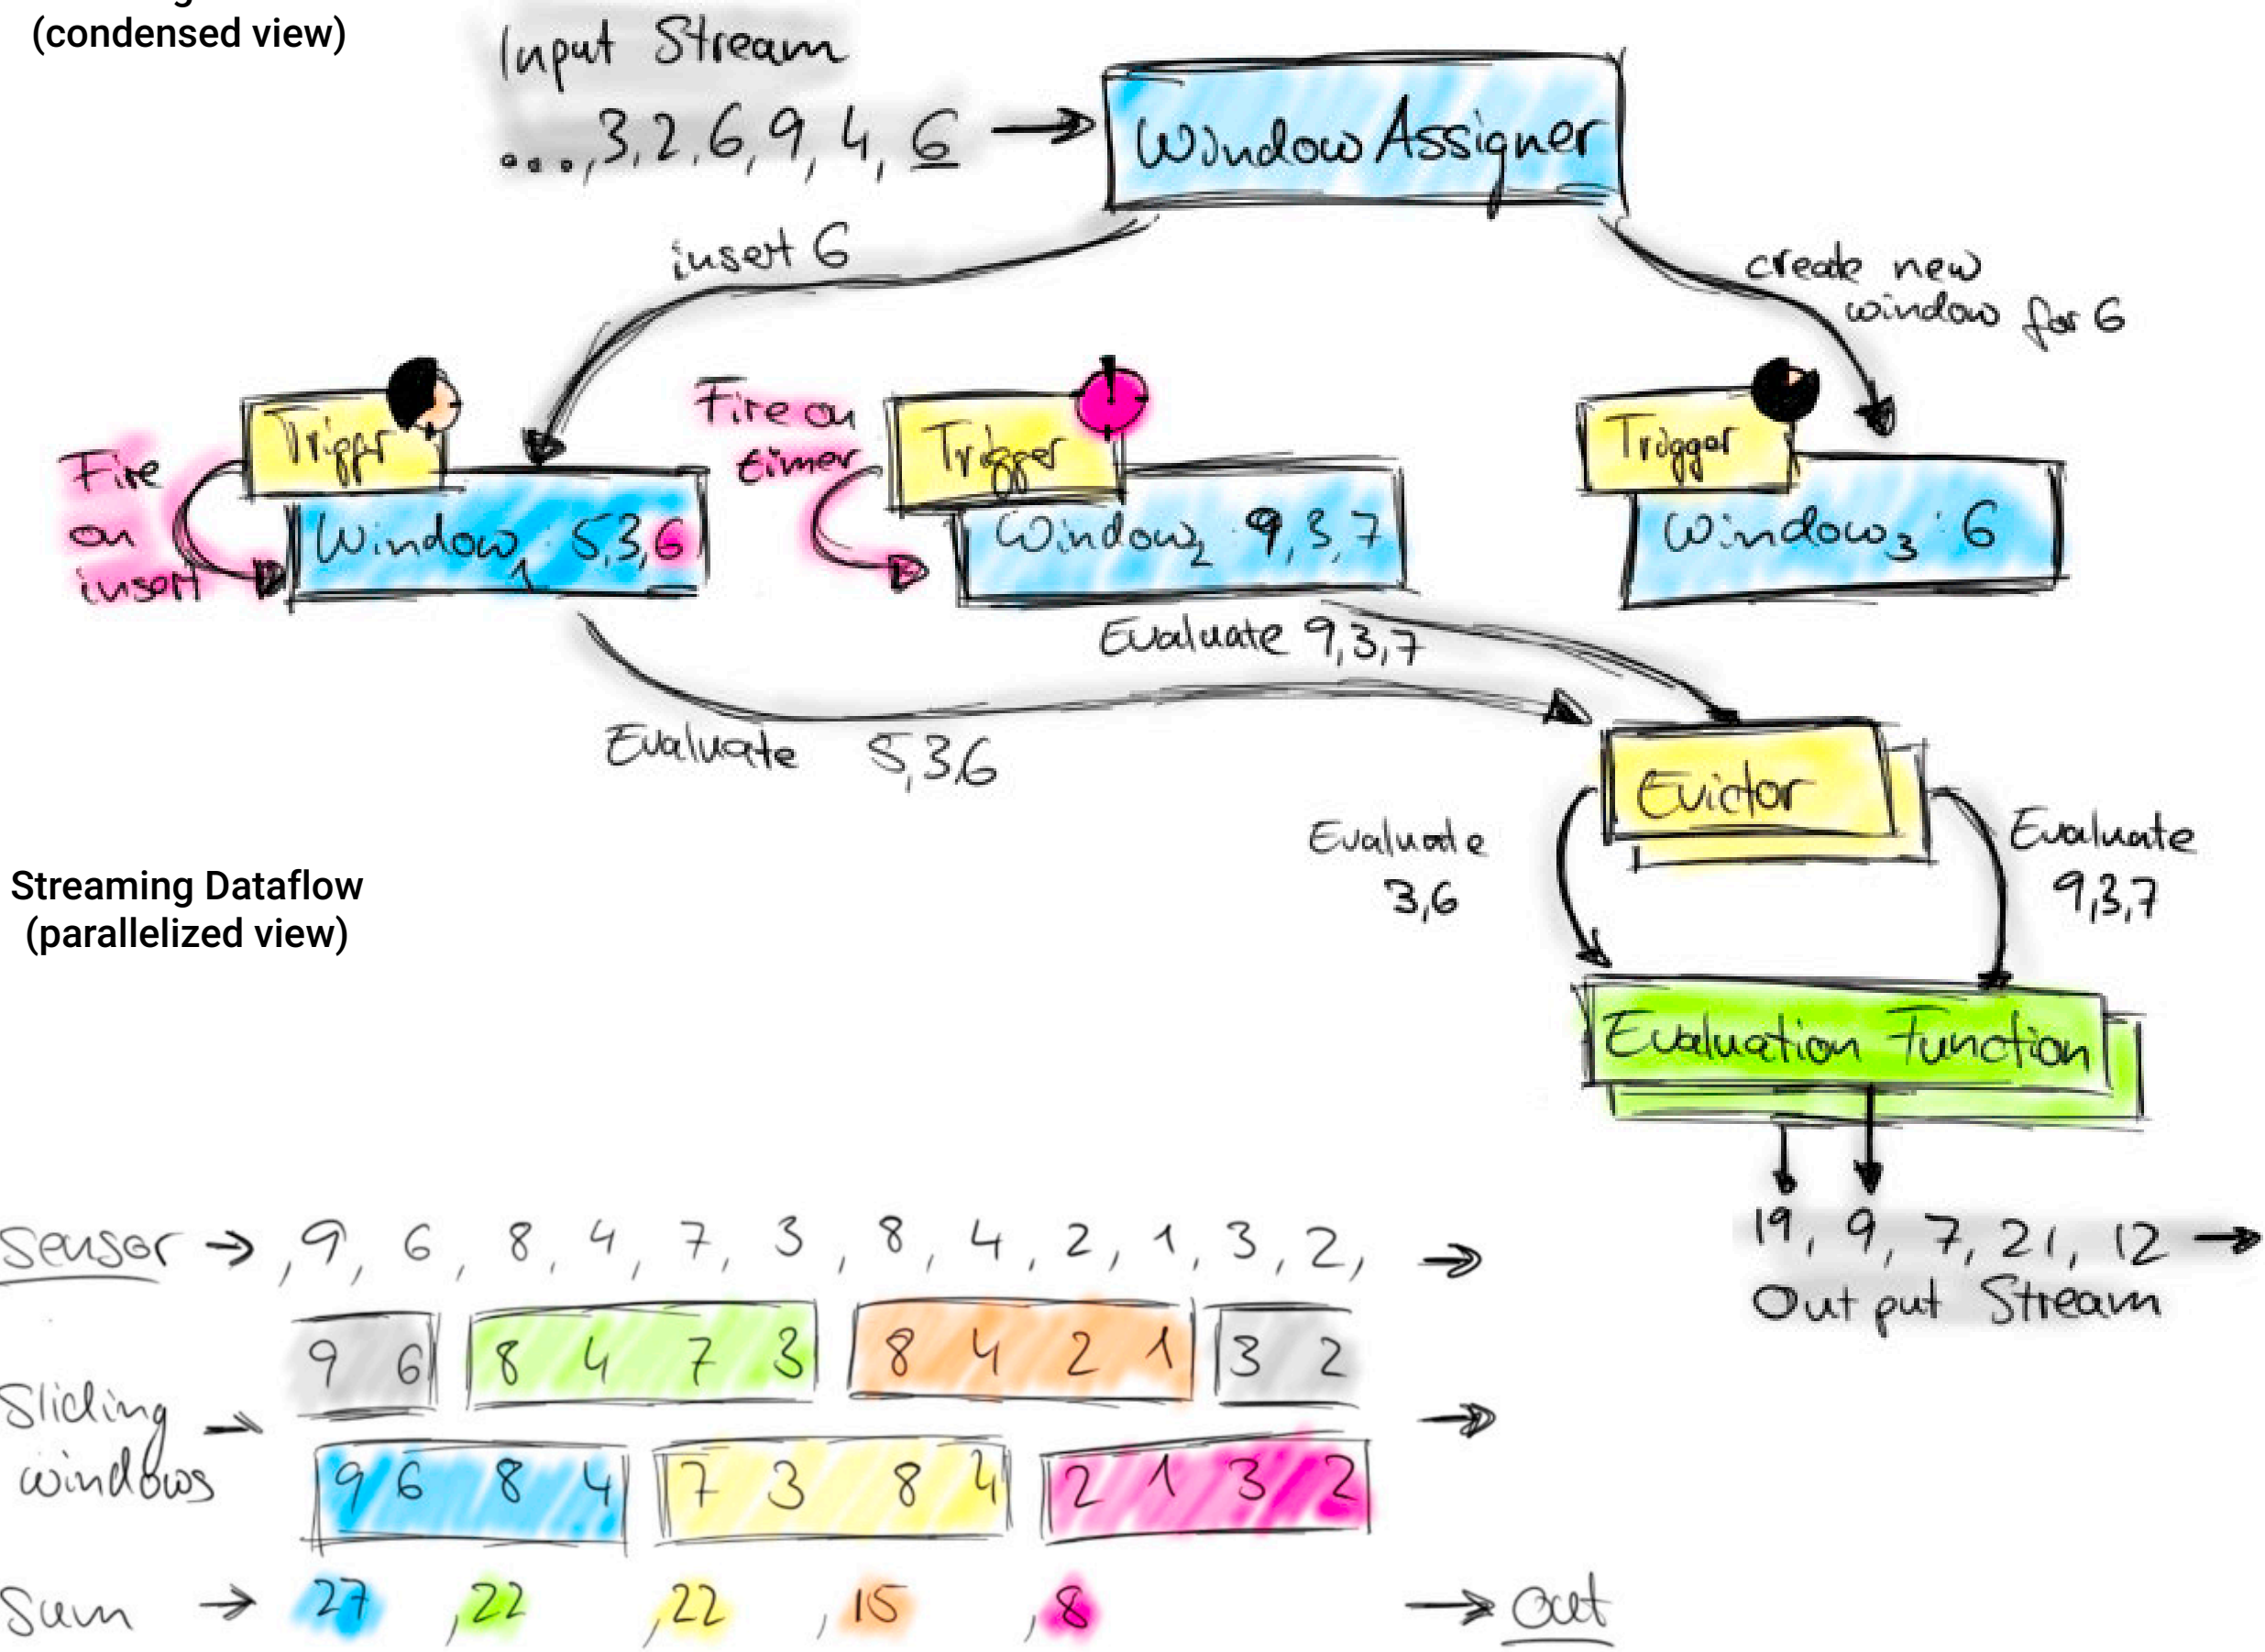


# Apache Flink: API II

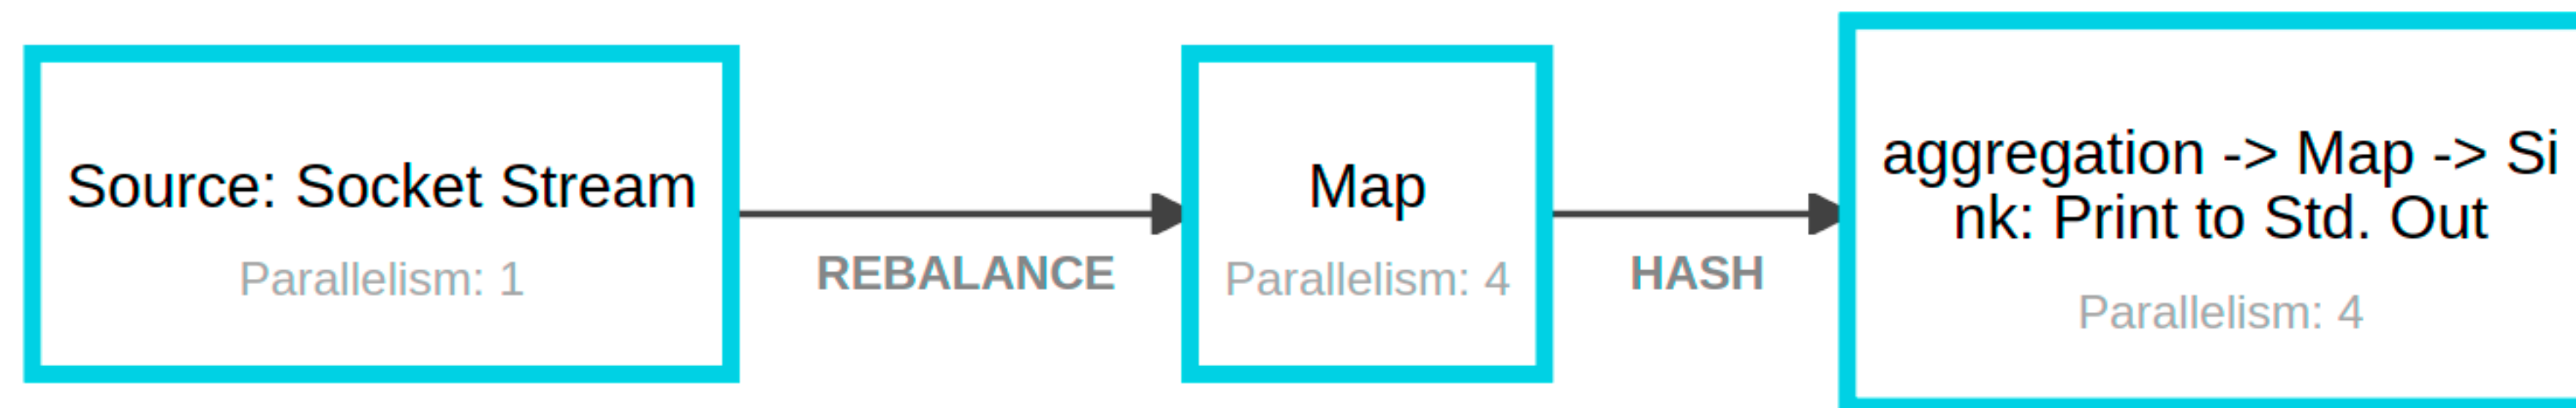


Streaming Dataflow (condensed view)

Streaming Dataflow (parallelized view)

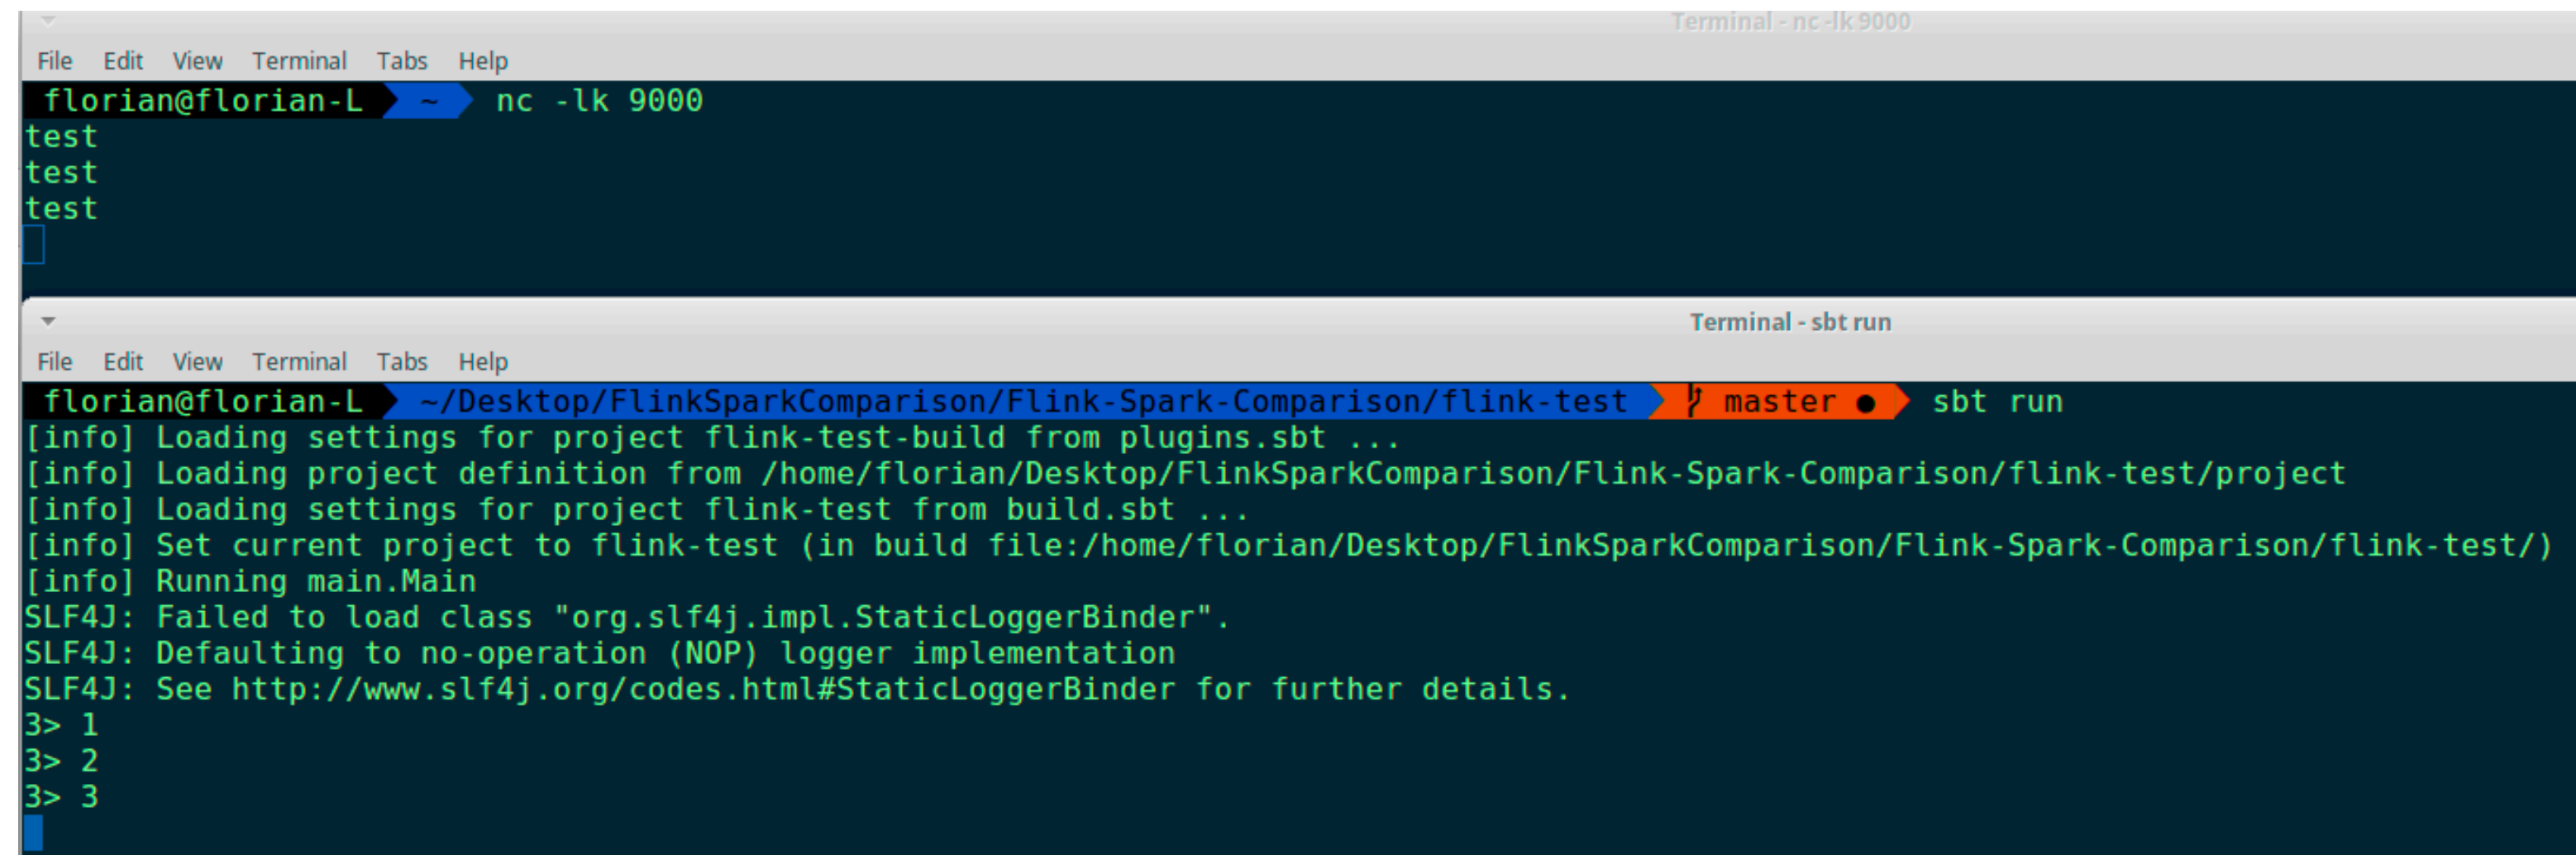


# Apache Flink: Streaming Example



```
val env: StreamExecutionEnvironment = StreamExecutionEnvironment.getExecutionEnvironment
env.setMaxParallelism(4)
val text: DataStream[String] = env.socketTextStream("127.0.0.1", 9000, '\n')
text.map(a=>(1,1))
    .keyBy(0)
    .sum(0)
    .map(b=>b._1)
    .print()
env.execute()
```

# Apache Flink: Streaming Example



The image shows two terminal windows. The top window, titled 'Terminal - nc -lk 9000', shows a netcat listener on port 9000 receiving three 'test' messages. The bottom window, titled 'Terminal - sbt run', shows the execution of an SBT project named 'flink-test'. It displays various log messages including project loading and SLF4J warnings, followed by three lines of output: '3> 1', '3> 2', and '3> 3'.

```
florian@florian-L ~ nc -lk 9000
test
test
test
[]

florian@florian-L ~/Desktop/FlinkSparkComparison/Flink-Spark-Comparison/flink-test master ● sbt run
[info] Loading settings for project flink-test-build from plugins.sbt ...
[info] Loading project definition from /home/florian/Desktop/FlinkSparkComparison/Flink-Spark-Comparison/flink-test/project
[info] Loading settings for project flink-test from build.sbt ...
[info] Set current project to flink-test (in build file:/home/florian/Desktop/FlinkSparkComparison/Flink-Spark-Comparison/flink-test/)
[info] Running main.Main
SLF4J: Failed to load class "org.slf4j.impl.StaticLoggerBinder".
SLF4J: Defaulting to no-operation (NOP) logger implementation
SLF4J: See http://www.slf4j.org/codes.html#StaticLoggerBinder for further details.
3> 1
3> 2
3> 3
```

- stateful behaviour
- job graph
  - define data source
  - transformations (keyBy: repartitions the given stream and ensures parallelism, sum: stateful continuous elements)
  - execute graph

# Apache Spark and Flink Benchmark

---

- **example should be understandable and implementable**
- **test algorithm is a simple word count over a given document**
- **lipsum dataset is used with 50 paragraphs**
  - replicated multiple times to reach 10, 100, 1000 and 2000MB
- **spark and flink have the same base configuration**
  - Spark is using a 2 second batch interval
  - we decided to disable windowing (2 second interval) in Flink because of no performance gain
- **only looking at the duration of the main job**
  - ignoring latency, throughput, backpressure

# Benchmark: Spark Implementation

---

```
def split(line: String): Array[String] =  
{  
    line.toLowerCase().replaceAll(",", " ").replaceAll(".", " ").split(" ")  
}
```

```
val sc = spark.sparkContext  
val ssc = new StreamingContext(sc, Seconds(2))  
var counts: RDD[(String, Long)] = sc.emptyRDD
```

```
val data = ssc.textFileStream("file://...")  
    .filter(!_isEmpty())  
    .flatMap(split(_))  
val words = data.map(f => (f, 1L)).reduceByKey((a, b) => a + b)  
words.foreachRDD(rdd => {  
    val merged = counts.union(rdd)  
    counts = merged.reduceByKey((a, b) => a + b)  
})
```

```
ssc.start()  
ssc.awaitTerminationOrTimeout(20000)  
ssc.stop(false)
```

# Benchmark: Flink Implementation

---

```
case class WordWithCount(word:String,count:Long)
```

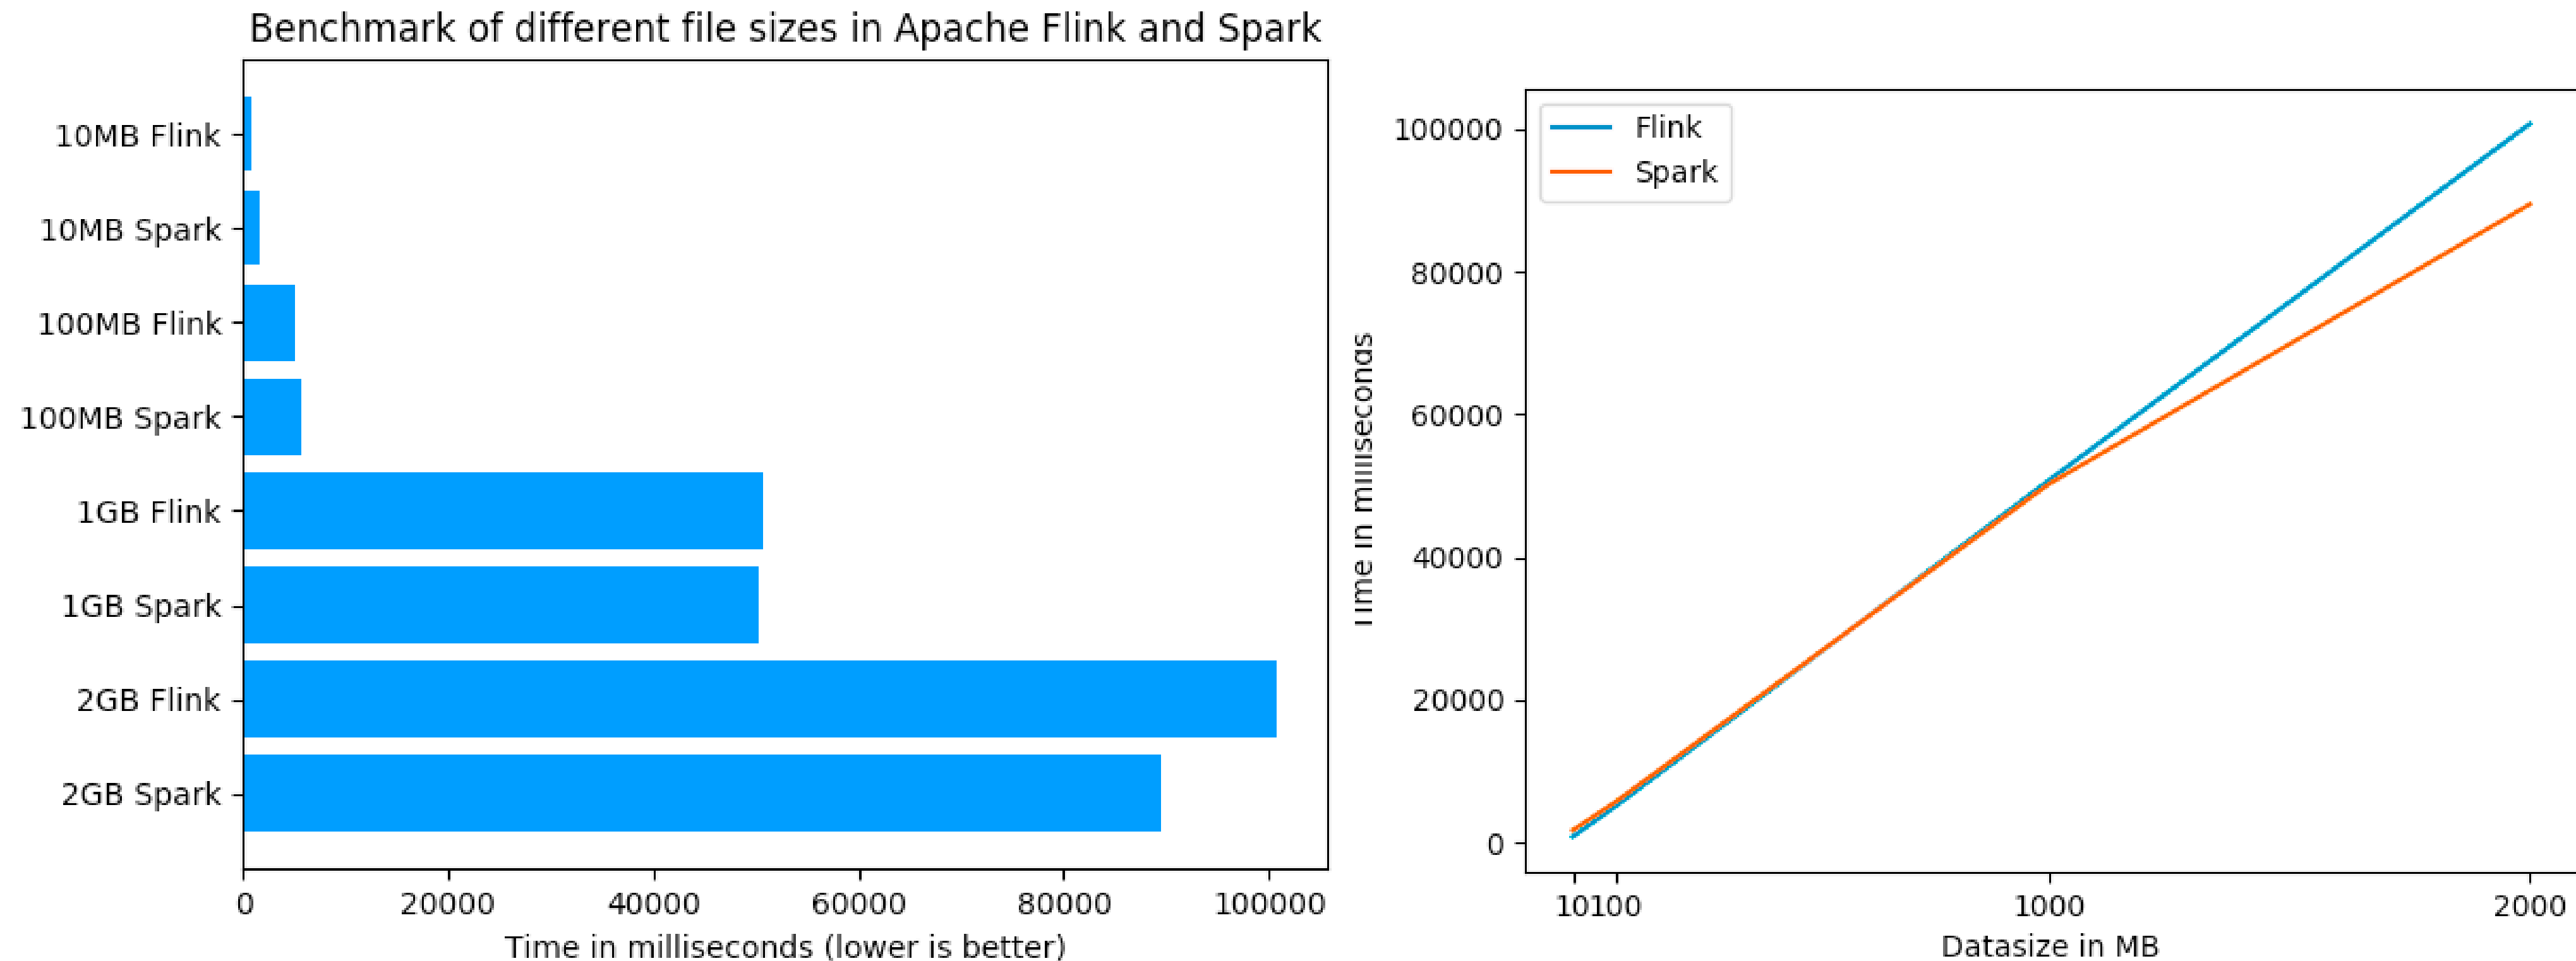
```
class LoremWordCount
```

```
{  
    val env: StreamExecutionEnvironment = StreamExecutionEnvironment.getExecutionEnvironment  
    env.setParallelism(4)  
    val text: DataStream[String] = env.readTextFile("...",StandardCharsets.UTF_8.name())  
    text.flatMap(line => line  
        .toLowerCase()  
        .replaceAll(",", " ")  
        .replaceAll(".", " ")  
        .split("\\s")  
    )  
    .map(w => WordWithCount(w, 1))  
    .keyBy("word").sum("count")  
    .addSink(new DiscardingSink[WordWithCount]())  
    env.execute()  
}
```

```
class Sink[T] extends SinkFunction[T] {}
```



# Benchmark: Results



- node specifications

- ubuntu 16.04.5 LTS

- Intel Core i5-7200U @ 2.50GHz (4 Threads), 8GB DDR4, 256 Samsung SSD



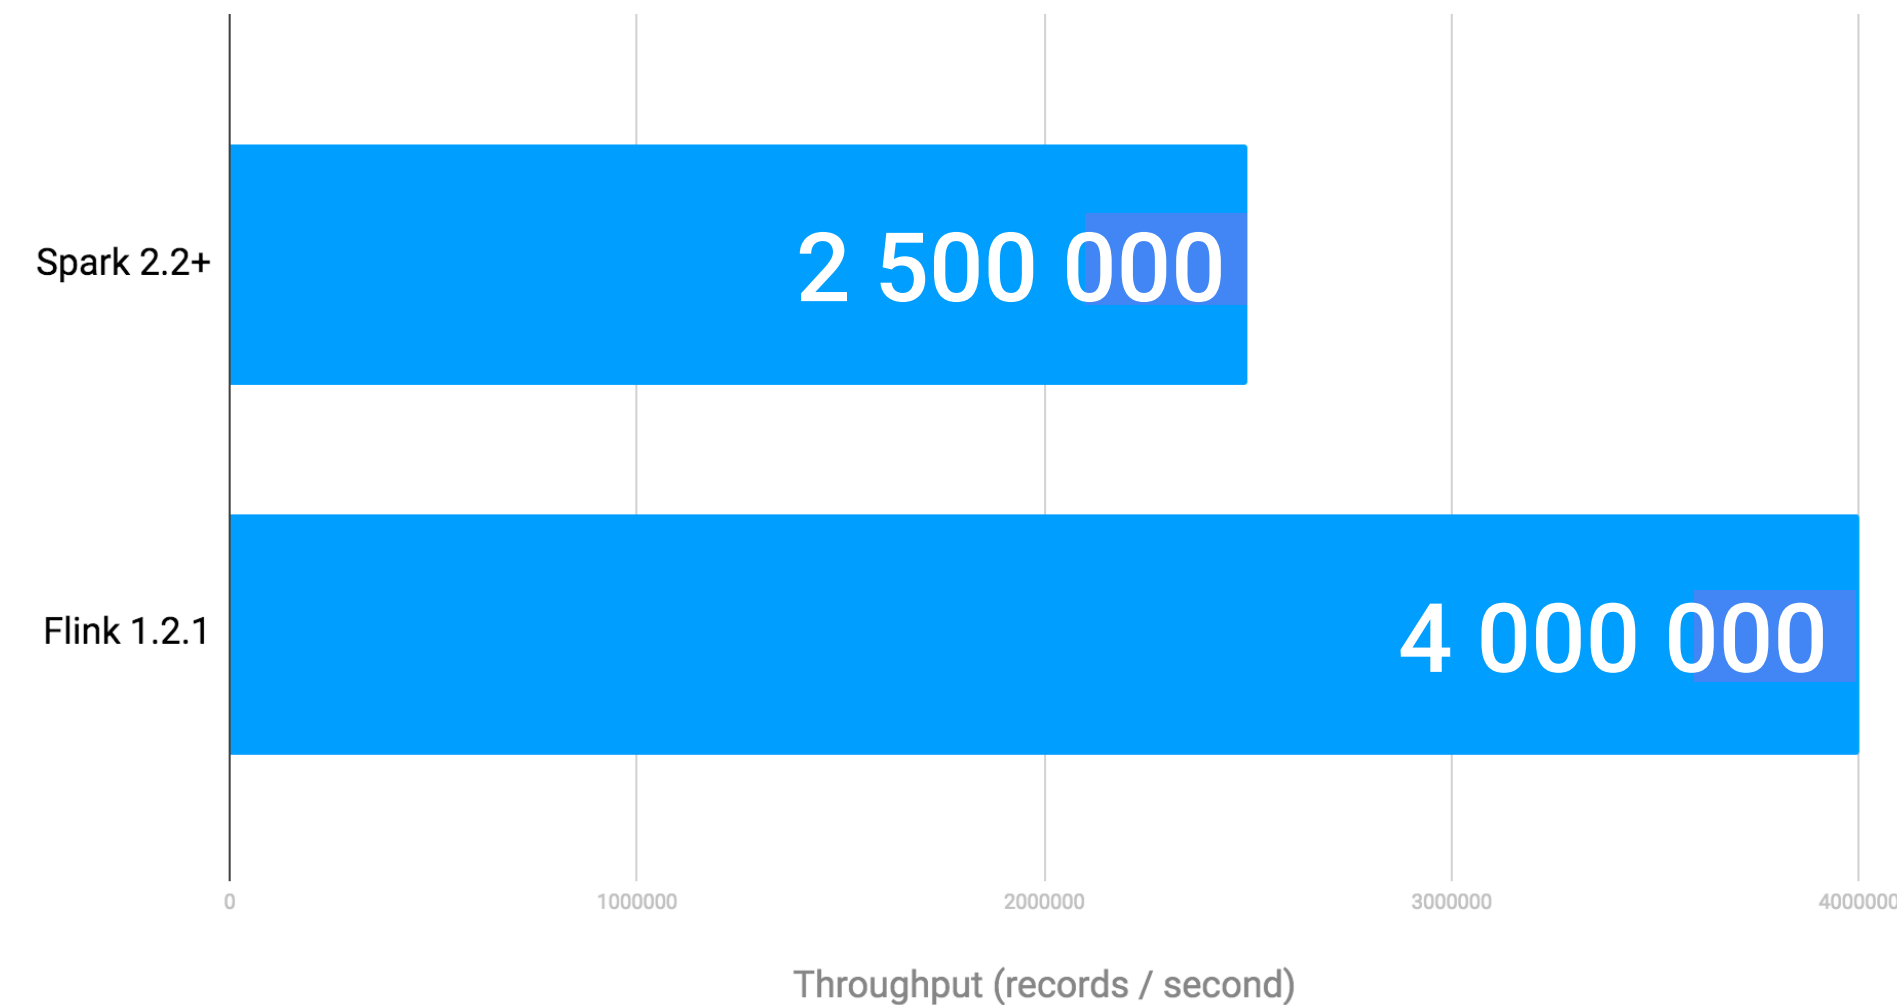
# Benchmark: Problems

---

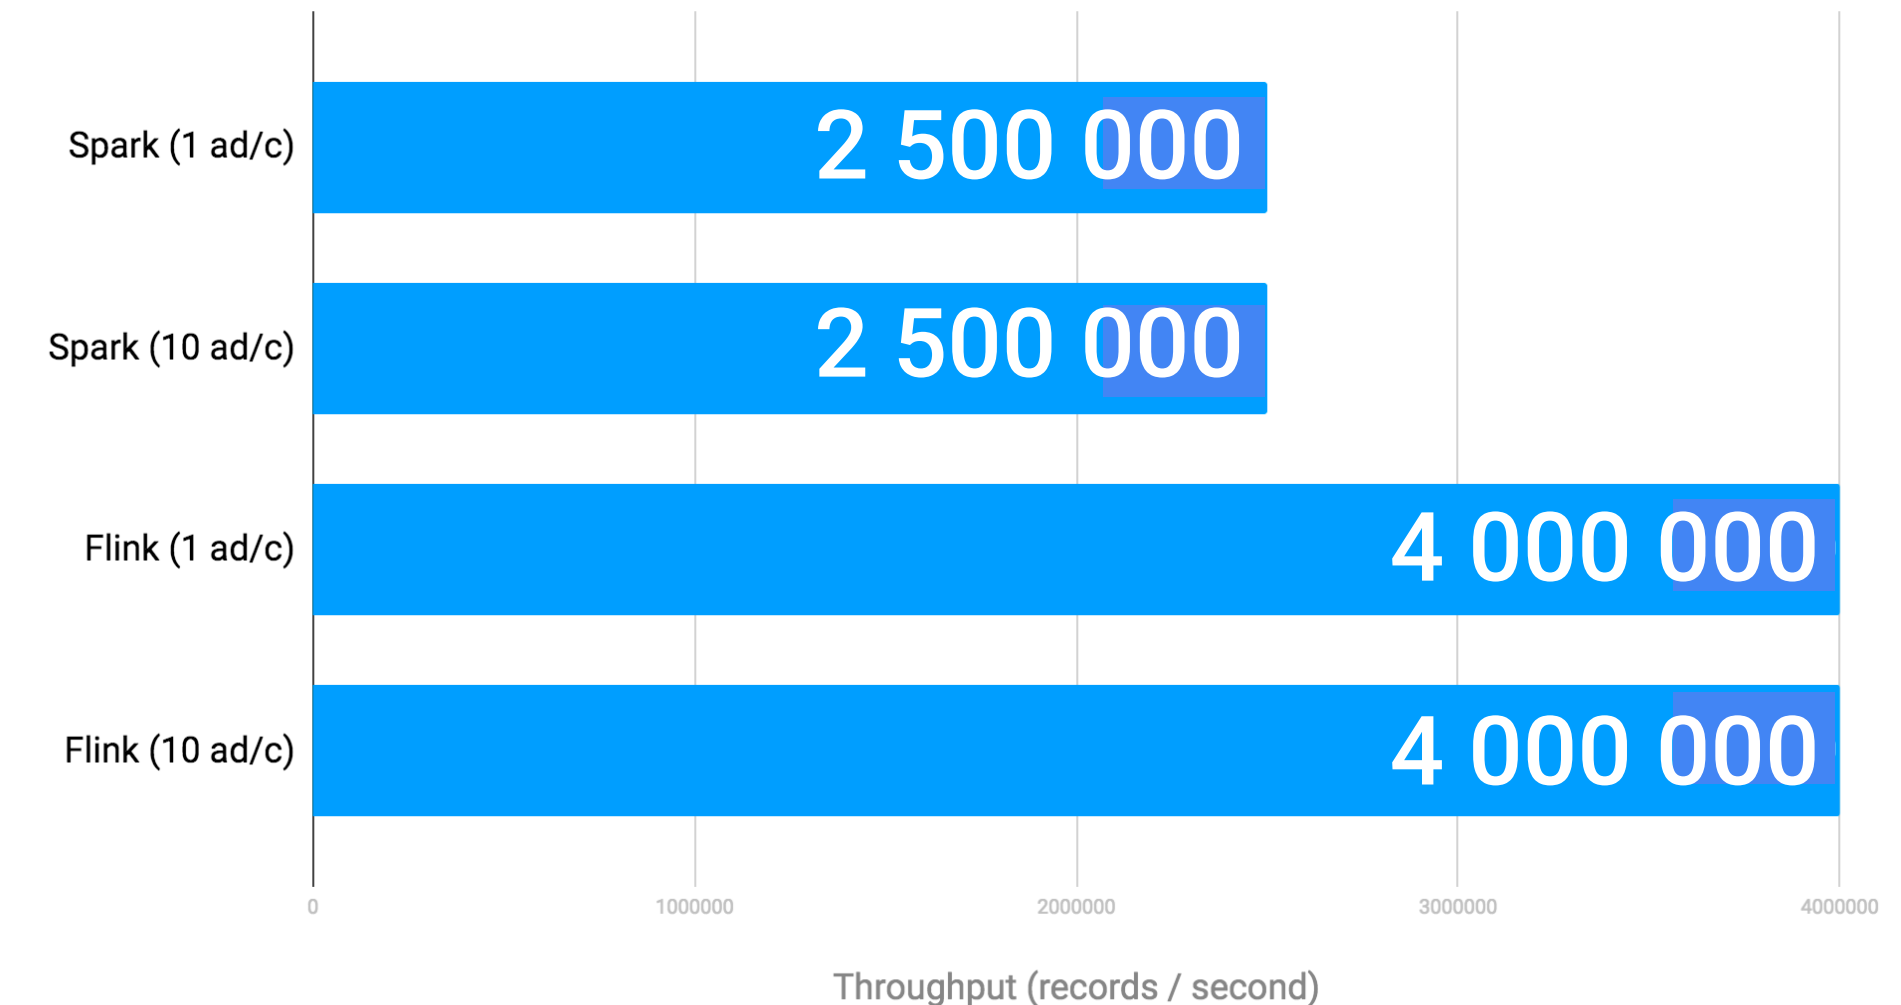
- **only a single data source**
  - reading from hard drive / ssd
- **the entry point is hard to define due to endless stream**
  - other metrics are needed to find the right comparisons
- **data will be discarded and not processed further**
- **one algorithm used**
  - different business logic has different performance requirements
- **local machine, scalability not tested**
  - maby CPU throttling and other issues

# Apache Spark and Flink other Benchmarks

Single Core Throughput (higher is better)



Single Core Throughput (higher is better)



»... that benchmark results most often represent a narrow combination of business logic and configuration options, deployed in an artificial environment.«

- data Artisans

# Sources

---

## Apache Spark

[spark.apache.org/streaming/](https://spark.apache.org/streaming/)  
[spark.apache.org/docs/2.2.0/streaming-programming-guide.html](https://spark.apache.org/docs/2.2.0/streaming-programming-guide.html)  
[spark.apache.org/docs/2.2.0/mllib-clustering.html](https://spark.apache.org/docs/2.2.0/mllib-clustering.html)  
[github.com/apache/spark/blob/master/mllib/src/main/scala/org/apache/spark/mllib/clustering/](https://github.com/apache/spark/blob/master/mllib/src/main/scala/org/apache/spark/mllib/clustering/)

## Apache Flink

“Benchmarking Distributed Stream Processing Engines”: Jeyhun Karimov, Tilmann Rabl, Asterios Katsifodimos, Roman Samarev, Henri Heiskanern, Volker Markl url: [arxiv.org/pdf/1802.08496.pdf](https://arxiv.org/pdf/1802.08496.pdf) [19/12/2018]  
[www.oreilly.com/library/view/stream-processing-with/9781491974285/](https://www.oreilly.com/library/view/stream-processing-with/9781491974285/)  
[www.kdnuggets.com/2016/10/beginners-guide-apache-flink-explained.html](https://www.kdnuggets.com/2016/10/beginners-guide-apache-flink-explained.html)  
[ci.apache.org/projects/flink/flink-docs-release-1.7/concepts/programming-model.html](https://ci.apache.org/projects/flink/flink-docs-release-1.7/concepts/programming-model.html)  
[flink.apache.org/news/2015/12/04/Introducing-windows.html](https://flink.apache.org/news/2015/12/04/Introducing-windows.html)  
[thirdeyeddata.io/apache-flink/](https://thirdeyeddata.io/apache-flink/)

## other Benchmarks

[www.data-artisans.com/blog/curious-case-broken-benchmark-revisiting-apache-flink-vs-databricks-runtime](https://www.data-artisans.com/blog/curious-case-broken-benchmark-revisiting-apache-flink-vs-databricks-runtime)  
[github.com/yahoo/streaming-benchmarks](https://github.com/yahoo/streaming-benchmarks)

## datasets

[exoplanetarchive.ipac.caltech.edu](https://exoplanetarchive.ipac.caltech.edu)  
[www.lipsum.com](https://www.lipsum.com)

# Appendix: Flink / Spark Code Comparision

---

```
object WordCount {  
  
  def main(args: Array[String]) {  
  
    val env = new SparkContext("local","wordCount")  
  
    val data = List("text1","text21 text22 text23","text3")  
  
    val dataRDD = env.parallelize(data)  
  
    val words = dataRDD.flatMap(value => value.split("\\s+"))  
  
    val mappedWords = words.map(value => (value,1))  
  
    val sum = mappedWords.reduceByKey(_+_)  
  
    println(sum.collect())  
  
  }  
}
```

```
object WordCount {  
  
  def main(args: Array[String]) {  
  
    val env = ExecutionEnvironment.getExecutionEnvironment  
  
    val data = List("text1","text21 text22 text23","text3")  
  
    val dataDataSet = env.fromCollection(data)  
  
    val words = dataDataSet.flatMap(value => value.split("\\s+"))  
  
    val mappedWords = words.map(value => (value,1))  
  
    val grouped = mappedWords.groupBy(0)  
  
    val sum = grouped.sum(1)  
  
    println(sum.collect())  
  }  
}
```